

COMP3211 Tutorial 2: Simple Agents

Fengming ZHU

Feb. 19&22, 2024

Department of CSE

HKUST

© 2024 Fengming Zhu. All rights reserved.

Overview

Production System

Boundary-Following Agents

Capabilities and Limitations

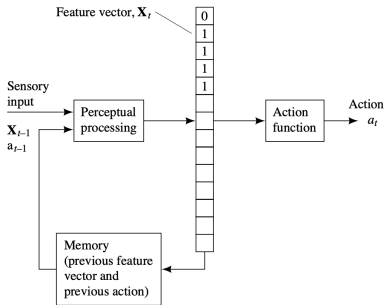
Genetic Programming

Biological description

Application in optimization

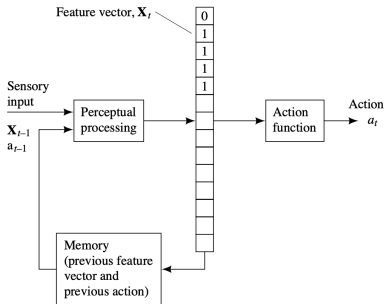
Overview

Simple Agents



© 1998 Morgan Kaufman Publishers

Simple Agents




© 1998 Morgan Kaufman Publishers

In terms of memory:

- Stimulus-response (reactive) agents,
- State machines,

Simple Agents



© 1998 Morgan Kaufman Publishers

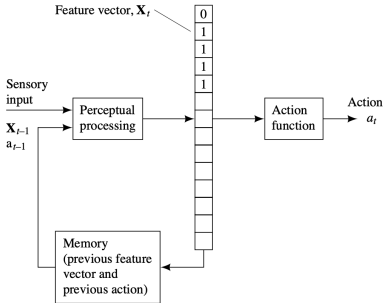
In terms of memory:

- Stimulus-response (reactive) agents,
- State machines,

In terms of action function:

- TLUs,
- Production systems

Simple Agents



© 1998 Morgan Kaufman Publishers

In terms of memory:

- Stimulus-response (reactive) agents,
- State machines,

In terms of action function:


- TLUs,
- Production systems

In terms of optimization:


- Error correction,
- Genetic programming

Production System


Boundary-Following Agents




x_1



x_2



x_3



x_4

$x_4\overline{x_1} \rightarrow \textit{north},$

$x_2\overline{x_3} \rightarrow \textit{south},$

$1 \rightarrow \textit{north}.$


$x_3\overline{x_4} \rightarrow \textit{west},$

$x_1\overline{x_2} \rightarrow \textit{east},$

Capabilities and Limitations

Example 1:


Can you find a production system by which the agent can reach the goal from any initial position?



Capabilities and Limitations

Example 2:


Can you find a production system by which the agent can reach the goal from any initial position?



Capabilities and Limitations

Example 2:


Can you find a production system by which the agent can reach the goal from any initial position?



How many past sensory readings should the agent remember?

Excercise

How about multi-agents:



Genetic Programming

Genetic process:

- Large enough population,

Genetic process:

- Large enough population,
- Survival of the fittest,

Genetic process:

- Large enough population,
- Survival of the fittest,
- Copy and Crossover,

Genetic process:

- Large enough population,
- Survival of the fittest,
- Copy and Crossover,
- Mutation

Example 3:

Find nice optima in the interval $[-1, 2]$ for the following function:


- $f(x) = x \sin(10\pi x) + 2$

Application in optimization

Example 3:

Find nice optima in the interval $[-1, 2]$ for the following function:

- $f(x) = x \sin(10\pi x) + 2$
- Quite complicated...



Example

Find nice optima in the interval $[-1, 2]$ for the following function:

- $f(x) = x \sin(10\pi x) + 2$

Example

Find nice optima in the interval $[-1, 2]$ for the following function:

- $f(x) = x \sin(10\pi x) + 2$

By genetic programming:

- Initialization: randomly generate “lots of”

$(x, f(x))$, representing (individual, fitness)

Example

Find nice optima in the interval $[-1, 2]$ for the following function:

- $f(x) = x \sin(10\pi x) + 2$

By genetic programming:

- Initialization: randomly generate “lots of”

$(x, f(x))$, representing (individual, fitness)

- Repeat for enough rounds:

Example

Find nice optima in the interval $[-1, 2]$ for the following function:

- $f(x) = x \sin(10\pi x) + 2$

By genetic programming:

- Initialization: randomly generate “lots of”

$(x, f(x))$, representing (individual, fitness)

- Repeat for enough rounds:

- Crossover: $x_{child} \leftarrow \lambda x_{father} + (1 - \lambda)x_{mother}$

Example

Find nice optima in the interval $[-1, 2]$ for the following function:

- $f(x) = x \sin(10\pi x) + 2$

By genetic programming:

- Initialization: randomly generate “lots of”

$(x, f(x))$, representing (individual, fitness)

- Repeat for enough rounds:
 - Crossover: $x_{child} \leftarrow \lambda x_{father} + (1 - \lambda)x_{mother}$
 - Mutation: $x \leftarrow \text{Normal}(x, 0.5)$, for 10% of individuals (x's)

Example

Find nice optima in the interval $[-1, 2]$ for the following function:

- $f(x) = x \sin(10\pi x) + 2$

By genetic programming:

- Initialization: randomly generate “lots of”

$(x, f(x))$, representing (individual, fitness)

- Repeat for enough rounds:
 - Crossover: $x_{child} \leftarrow \lambda x_{father} + (1 - \lambda)x_{mother}$
 - Mutation: $x \leftarrow \text{Normal}(x, 0.5)$, for 10% of individuals (x's)
 - Selection: choose top 2/3 fittest individuals (x's)

Example




Figure 1: Larger population and more generations

Compared to GD

In terms of optimization, why we still need gradient descent to train a TLU/neural network?

- Genetic programming: zero-order information
- GD (error-correction as a special case): first-order derivatives
- Newton's method: second order derivatives (computing inverse of Hessian matrix is hard)

Thanks!