

COMP3211 Tutorial 2: Simple Agents

Fengming ZHU

Feb. 19&22, 2024

Department of CSE
HKUST

Overview

Production System

Boundary-Following Agents

Capabilities and Limitations

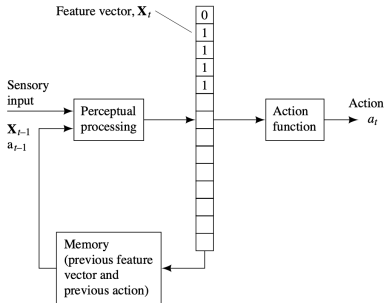
Genetic Programming

Biological description

Application in optimization

Overview

Simple Agents

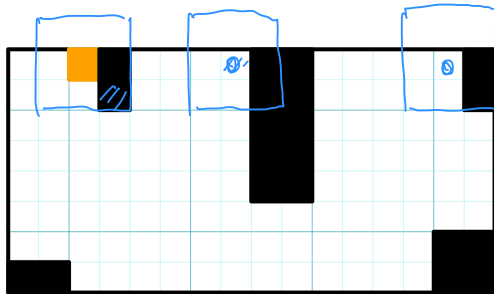


© 1998 Morgan Kaufman Publishers

Capabilities and Limitations

Example 2:

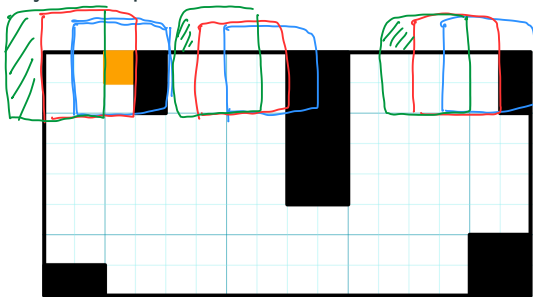
Can you find a production system by which the agent can reach the goal from any initial position?



Capabilities and Limitations

Example 2:

Can you find a production system by which the agent can reach the goal from any initial position?

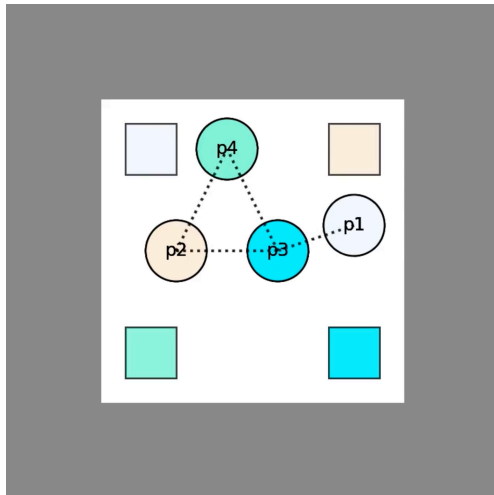


How many past sensory readings should the agent remember?

2

Excercise

How about multi-agents:



Genetic Programming

Genetic process:

- Large enough population,

Genetic process:

- Large enough population,
- Survival of the fittest,

Genetic process:

- Large enough population,
- Survival of the fittest,
- Copy and Crossover,

Genetic process:

- Large enough population,
- Survival of the fittest,
- Copy and Crossover,
- Mutation

Application in optimization

Example 3:

Find nice optima in the interval $[-1, 2]$ for the following function:

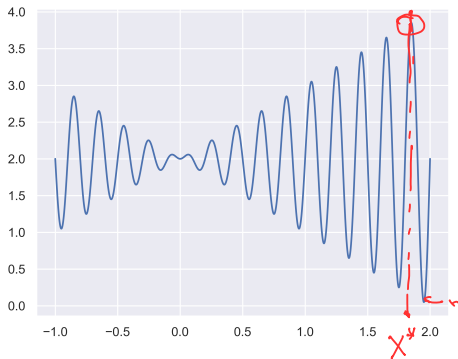
- $f(x) = x \sin(10\pi x) + 2$

Application in optimization

Example 3:

Find nice optima in the interval $[-1, 2]$ for the following function:

- $f(x) = x \sin(10\pi x) + 2$
- Quite complicated...



Example

Find nice optima in the interval $[-1, 2]$ for the following function:

- $f(x) = x \sin(10\pi x) + 2$

Example

Find nice optima in the interval $[-1, 2]$ for the following function:

- $f(x) = x \sin(10\pi x) + 2$

By genetic programming:

- Initialization: randomly generate “lots of”

$(x, f(x))$, representing (individual, fitness)

Example

Find nice optima in the interval $[-1, 2]$ for the following function:

- $f(x) = x \sin(10\pi x) + 2$

By genetic programming:

- Initialization: randomly generate “lots of”

$(x, f(x))$, representing (individual, fitness)

- Repeat for enough rounds:

Example

Find nice optima in the interval $[-1, 2]$ for the following function:

- $f(x) = x \sin(10\pi x) + 2$

By genetic programming:

- Initialization: randomly generate “lots of”

$(x, f(x))$, representing (individual, fitness)

- Repeat for enough rounds:

- Crossover: $\underline{x_{child}} \leftarrow \lambda \underline{x_{father}} + (1 - \lambda) \underline{x_{mother}}$

Example

Find nice optima in the interval $[-1, 2]$ for the following function:

- $f(x) = x \sin(10\pi x) + 2$

By genetic programming:

- Initialization: randomly generate “lots of”

$(x, f(x))$, representing (individual, fitness)

- Repeat for enough rounds:

- Crossover: $x_{child} \leftarrow \lambda x_{father} + (1 - \lambda)x_{mother}$

- Mutation: $x \leftarrow \text{Normal}(x, 0.5)$, for 10% of individuals (x's)



Example

Find nice optima in the interval $[-1, 2]$ for the following function:

- $f(x) = x \sin(10\pi x) + 2$

By genetic programming:

- Initialization: randomly generate “lots of”

$(x, f(x))$, representing (individual, fitness)

- Repeat for enough rounds:

← • Crossover: $x_{child} \leftarrow \lambda x_{father} + (1 - \lambda)x_{mother}$

• Mutation: $x \leftarrow \text{Normal}(x, 0.5)$, for 10% of individuals (x's)

← • Selection: choose top 2/3 fittest individuals (x's)

Example

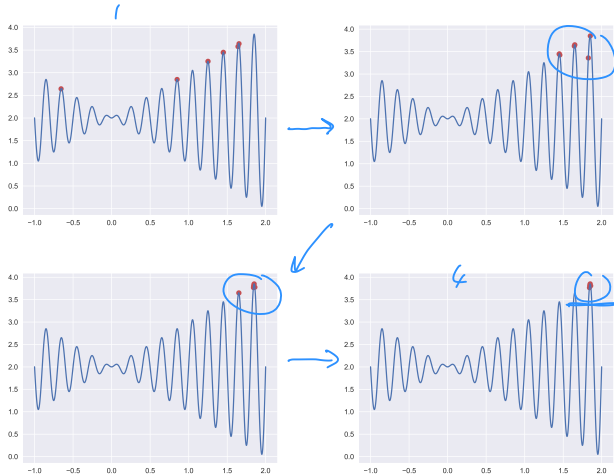
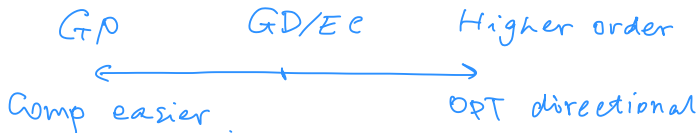


Figure 1: Larger population and more generations

Compared to GD

In terms of optimization, why we still need gradient descent to train a TLU/neural network?

- Genetic programming: zero-order information
- GD (error-correction as a special case): first-order derivatives
- Newton's method: second order derivatives (computing inverse of Hessian matrix is hard)



Thanks!