

COMP3211 Tutorial 3: Search

Fengming ZHU

Feb. 26&29, 2024

Department of CSE

HKUST

© 2024 Fengming Zhu. All rights reserved.

Search

Before Formulations

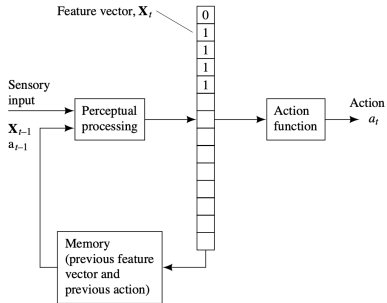
Formulation

Search Diagram

Exercise

Search

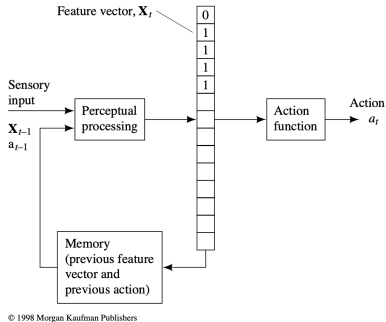
Simple Agents



© 1998 Morgan Kaufman Publishers

Figure 1: Simple agents

Simple Agents



Key points:

- Respond to the environment,

Figure 1: Simple agents

Simple Agents




Figure 1: Simple agents

Key points:

- Respond to the environment,
- Consider how the world IS, or HAVE BEEN,

Simple Agents

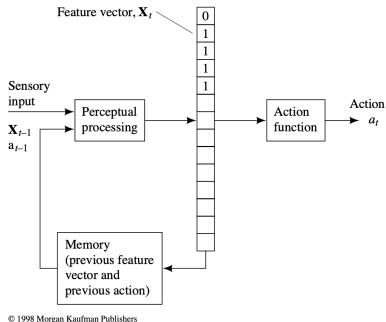


Figure 1: Simple agents

Key points:

- Respond to the environment,
- Consider how the world IS, or HAVE BEEN,
- Cannot imagine how the world WOULD BE.

Notations:

- A set of states \mathcal{S}
- An initial state $I \in \mathcal{S}$
- A goal state $G \in \mathcal{S}$ (sometimes a goal test)
- A set of actions \mathcal{A}
- Deterministic transitions $T : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$
- Cost function $c : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$
- A solution (path) is a sequence of actions from I to G .

Follow-up Questions

Question #1:

Can reactive agents or state machines do search?

Follow-up Questions

Question #1:

Can reactive agents or state machines do search?

- Neither, should be agents that can plan ahead (transitions matter).

Follow-up Questions

Question #1:

Can reactive agents or state machines do search?

- Neither, should be agents that can plan ahead (transitions matter).

Question #2:

For agent 1, who can compute a feasible plan (path), if you extract her plan and deploy to agent 2 under the same setting, who has no sensing ability and no computing power, can she successfully reach the goal?

Follow-up Questions

Question #1:

Can reactive agents or state machines do search?



- Neither, should be agents that can plan ahead (transitions matter).

Question #2:

For agent 1, who can compute a feasible plan (path), if you extract her plan and deploy to agent 2 under the same setting, who has no sensing ability and no computing power, can she successfully reach the goal?

- Yes, once computed, just blindly execute it.

General Idea: Graph/Tree Search



A node can only be expanded once, while it may be visited multiple times!


Key concepts:

- Fringe (frontier)
- Expansion
- Exploration strategy

Breadth-First Search

Strategy: expand a shallowest node first

Implementation: Fringe is a FIFO queue



Breadth-First Search

Strategy: expand a shallowest node first

Implementation: Fringe is a FIFO queue






Figure 2: Breadth-First Search

Uniform Cost Search

*Strategy: expand a
cheapest node first:*

*Fringe is a priority queue
(priority: cumulative cost)*



Uniform Cost Search

*Strategy: expand a
cheapest node first:*

*Fringe is a priority queue
(priority: cumulative cost)*





Figure 3: Uniform Cost Search

Greedy Search

- Strategy: expand a node that you think is closest to a goal state.
- Best case: every time you make a perfect guess.
- Worst case: turn around until you get into a dead end.





Figure 4: Uniform Cost Search

- Uniform-cost orders by path cost, or backward cost $g(n)$
- Greedy orders by goal proximity, or forward cost $h(n)$

- Uniform-cost orders by path cost, or backward cost $g(n)$
- Greedy orders by goal proximity, or forward cost $h(n)$
- Strategy: expand a node that is best so far.

A* Search

- Uniform-cost orders by path cost, or backward cost $g(n)$
- Greedy orders by goal proximity, or forward cost $h(n)$
- Strategy: expand a node that is best so far.
- Implementation: priority queue, $f(n) = g(n) + h(n)$.

- Uniform-cost orders by path cost, or backward cost $g(n)$
- Greedy orders by goal proximity, or forward cost $h(n)$
- Strategy: expand a node that is best so far.
- Implementation: priority queue, $f(n) = g(n) + h(n)$.
- Admissible heuristic: $h(n) \leq \text{cost}(n, G)$.

A* Search

- Uniform-cost orders by path cost, or backward cost $g(n)$
- Greedy orders by goal proximity, or forward cost $h(n)$
- Strategy: expand a node that is best so far.
- Implementation: priority queue, $f(n) = g(n) + h(n)$.
- Admissible heuristic: $h(n) \leq \text{cost}(n, G)$.




Figure 5: $h(n) = \text{shortest_path_length}(n, G)$






Figure 6:

$h(n) = \text{shortest_path_length}(n, G)$

A* Search – another example



- Uniform-cost orders by path cost, or backward cost $g(n)$
- Greedy orders by goal proximity, or forward cost $h(n)$
- Strategy: expand a node that is best so far.
- Implementation: priority queue, $f(n) = g(n) + h(n)$.
- Admissable heuristic: $h(n) \leq \text{cost}(n, G)$.
- Live demo: <https://www.movingai.com/SAS/index.html>

A^* search:

- An offline process – when the map changes: need to replan the whole solution.

A^* search:

- An offline process – when the map changes: need to replan the whole solution.
- Can we reuse any historical data to be more efficient?


A^* search:

- An offline process – when the map changes: need to replan the whole solution.
- Can we reuse any historical data to be more efficient?

Exercise

A* search:

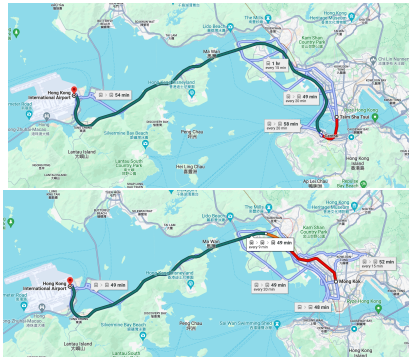
- An offline process – when the map changes: need to replan the whole solution.
- Can we reuse any historical data to be more efficient?



Exercise

A* search:

- An offline process – when the map changes: need to replan the whole solution.
- Can we reuse any historical data to be more efficient?



Exercise

A^* search:

- An offline process – when the map changes: need to replan the whole solution.
- Can we reuse any historical data to be more efficient?




Figure 7: Plan backwards from the goal to the start: D^* Lite

Thanks!